Tutorial 7

Apache Cassandra

To get started with Cassandra NoSQL database, we will step through a single-node, local installation on VM.

1) The following points are the requirements to run Cassandra locally: Move to the Apache download site for the Cassandra project (http://cassandra.apache.org/download/), choose 3.11.16, and select a mirror to download the latest version of Cassandra. When complete, copy the .tar or .gzip file to a location that your user has read and write permissions for. This example will assume that this is going to be the ~/Downloads/ directory on ubuntu VM.

Note: If you could not understand from the command instructions, then please check the screenshot for better understanding. If you like to explore the details of each command along with examples, then check the documentation of Cassandara on website: https://cassandra.apache.org/doc/latest/

Download Apache Cassandra from the following link as mentioned below

https://dlcdn.apache.org/cassandra/3.11.17/apache-cassandra-3.11.17-bin.tar.gz

2) `$cd Downloads`

Follow the instructions to unzip on the below screenshots and change the name of the folder as you did during the Hadoop, HBase and Spark installations (Tutorials, 2, 4, 6).

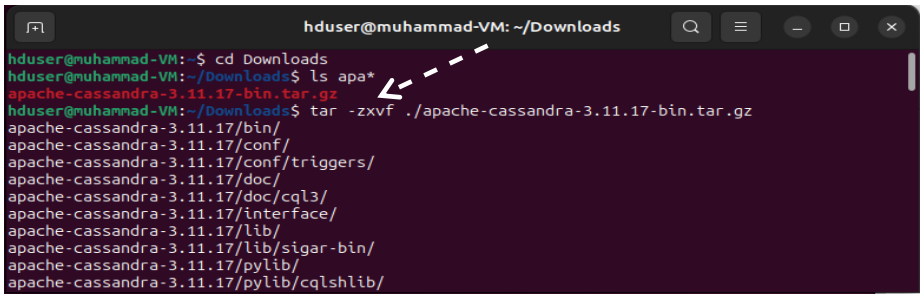

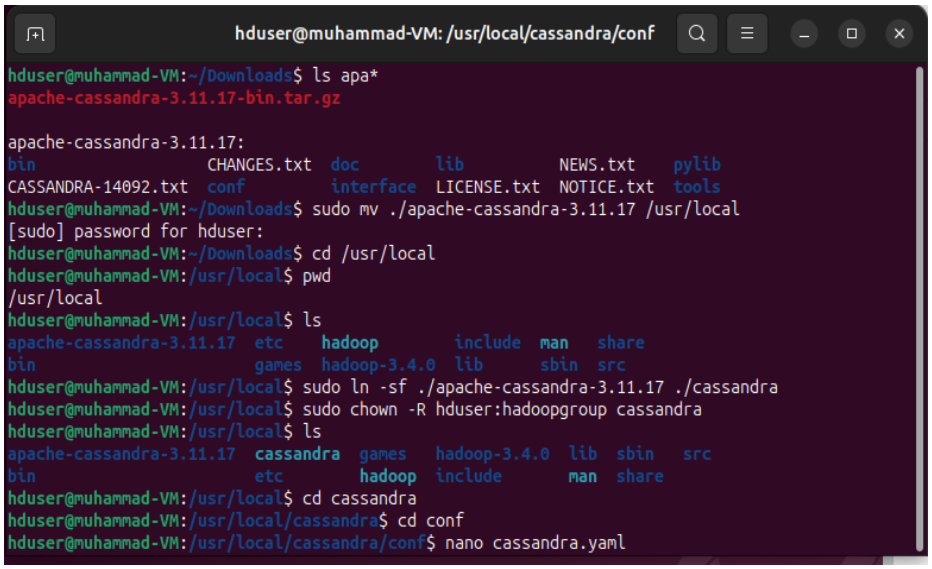

3) Configuration: At this point, you could start your node with no further configuration. However, it is good to get into the habit of checking and adjusting the properties that are indicated as follows using instructions as shown in the screenshot in step no. 2.

```
$cd /usr/local
$cd cassandra
$cd conf
$nano cassandra.yaml
```

It is usually a good idea to rename your cluster. Inside the conf/cassandra.yaml file, specify a new cluster_name property, overwriting the default Test Cluster as shown below in screenshot:

```
cluster_name: 'BDSP Cluster'
```

The num_tokens property default of 256 has proven to be too high for the newer, 3.x versions of Cassandra. Go ahead and set that to 24:

```
num_tokens: 24
```

save the file using nano/ gedit editor after updating above parameters.

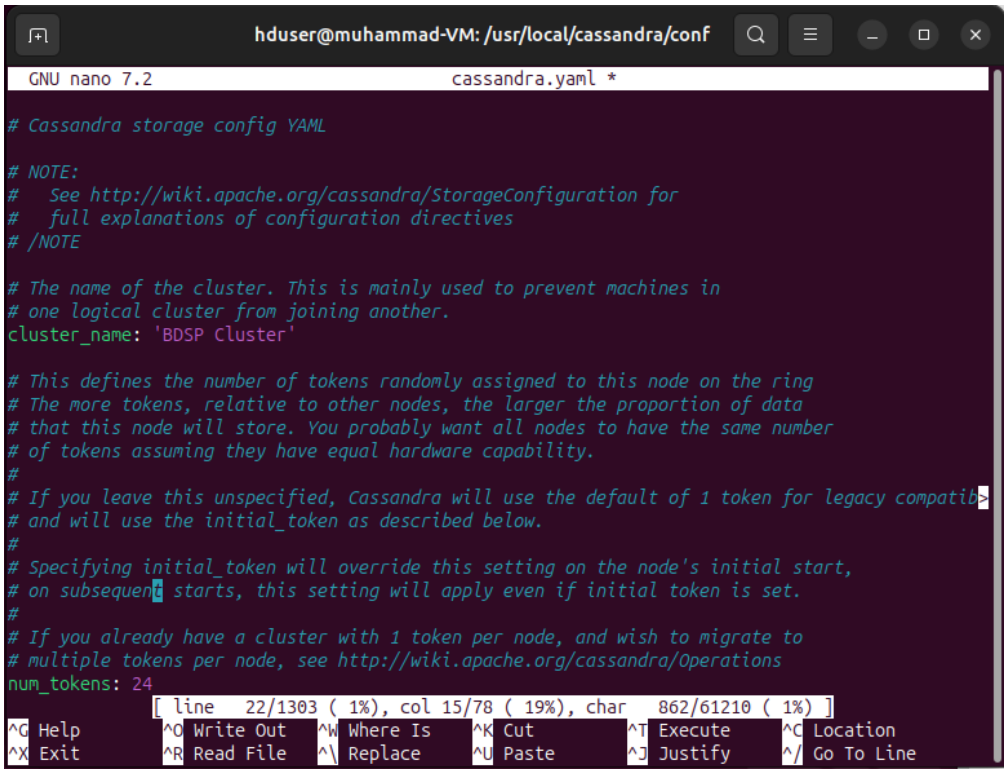

Press Alt+C to display the line number using nano editor.

4) By default, Cassandra will come up bound to localhost or 127.0.0.1. For your own local development machine

```
hduser@muhammad-VM:/usr/local/cassandra/conf$ nano cassandra-rackdc.properties
```

```
cassandra-rackdc.properties
```

In terms of NoSQL databases, Apache Cassandra handles multi-data center awareness better than any other. To configure this, each node must use **GossipingPropertyFileSnitch** (as previously mentioned in the preceding **cassandra.yaml** configuration process) and must have its local data center (and rack) settings defined. Therefore, we set the dc and rack properties in the **conf/cassandra-rackdc.properties** file:

```
dc=ClockworkAngels
rack=R40
```

If these properties are already set as shown below in the screenshot. You can exit from the editor.

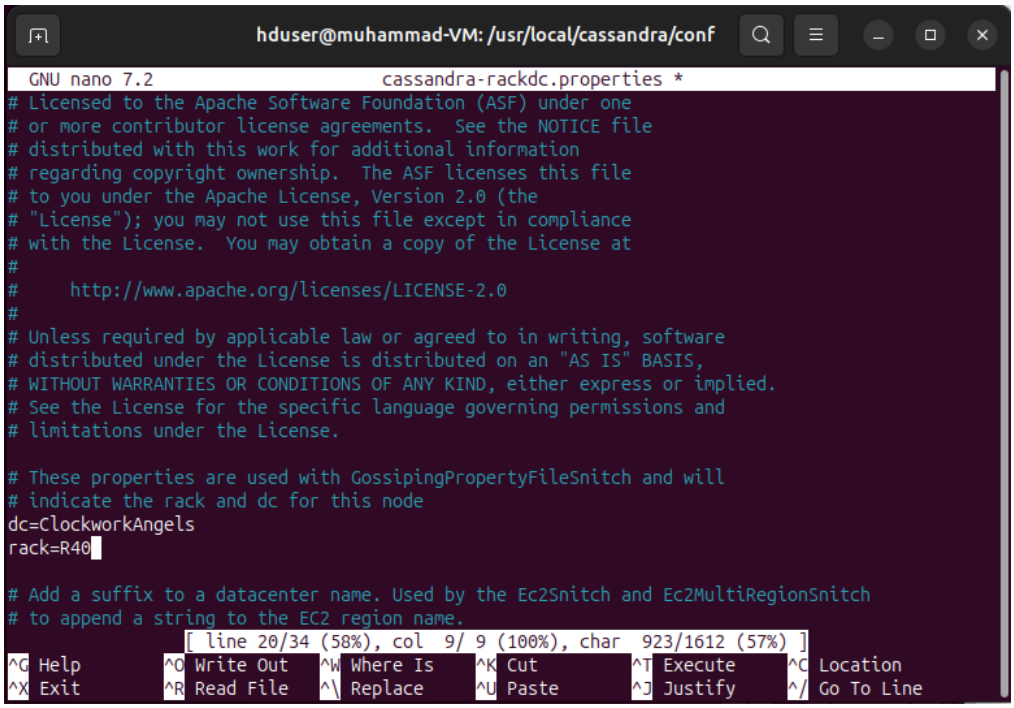

5) Starting Cassandra: To start Cassandra locally, execute the Cassandra script. If no arguments are passed, it will run in the foreground. To have it run in the background, send the **-p flag** with a destination file for the **Process ID** (**PID**):

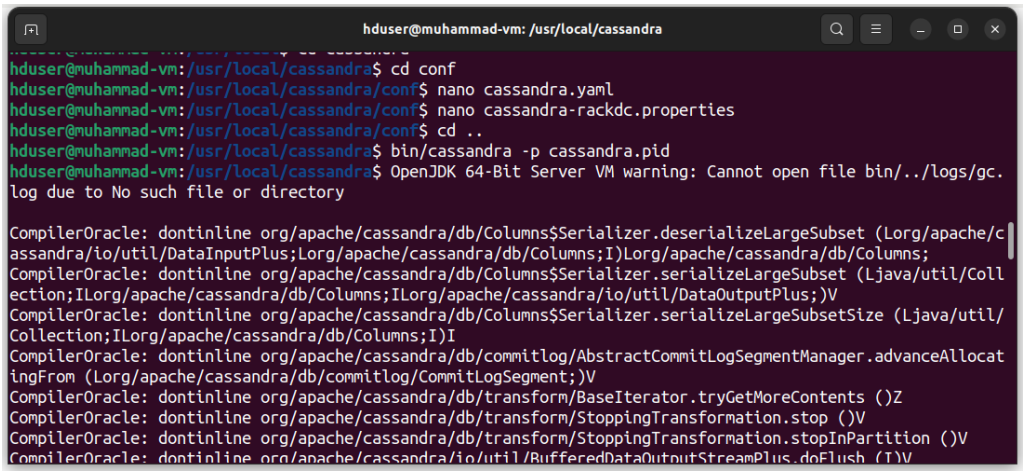


```
$bin/cassandra -p cassandra.pid
```

Or

```
$bin/cassandra -f
```

Now the Cassandra started. When you press the Enter key, Cassandra stopped automatically first time only as shown in the below screenshot.

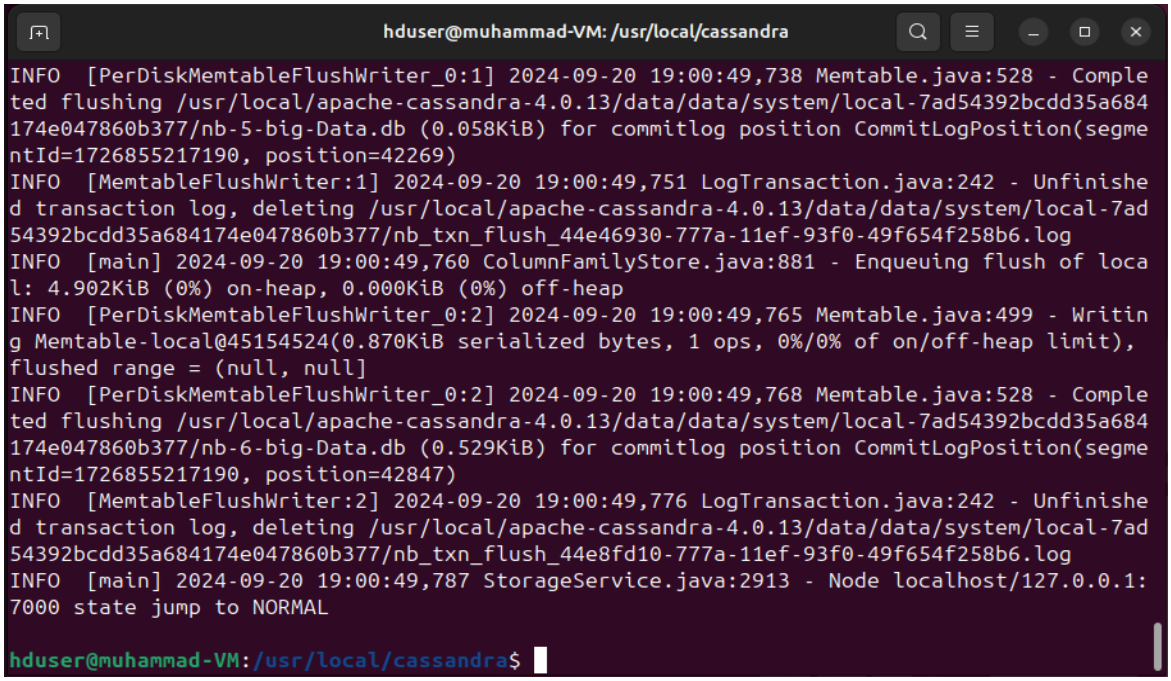

Keep this window open for now; please don't close it. However, if Cassandra is working in the first time, press Ctrl + C to stop the Cassandra process to avoid the errors in the future. You can use the command to kill the process as mentioned below on the screen. **You must perform these steps if the Cassandra showed an error in the case of start-up using this command (`bin/cassandra -p cassandra.pid` or `bin/cassandra -f`).**

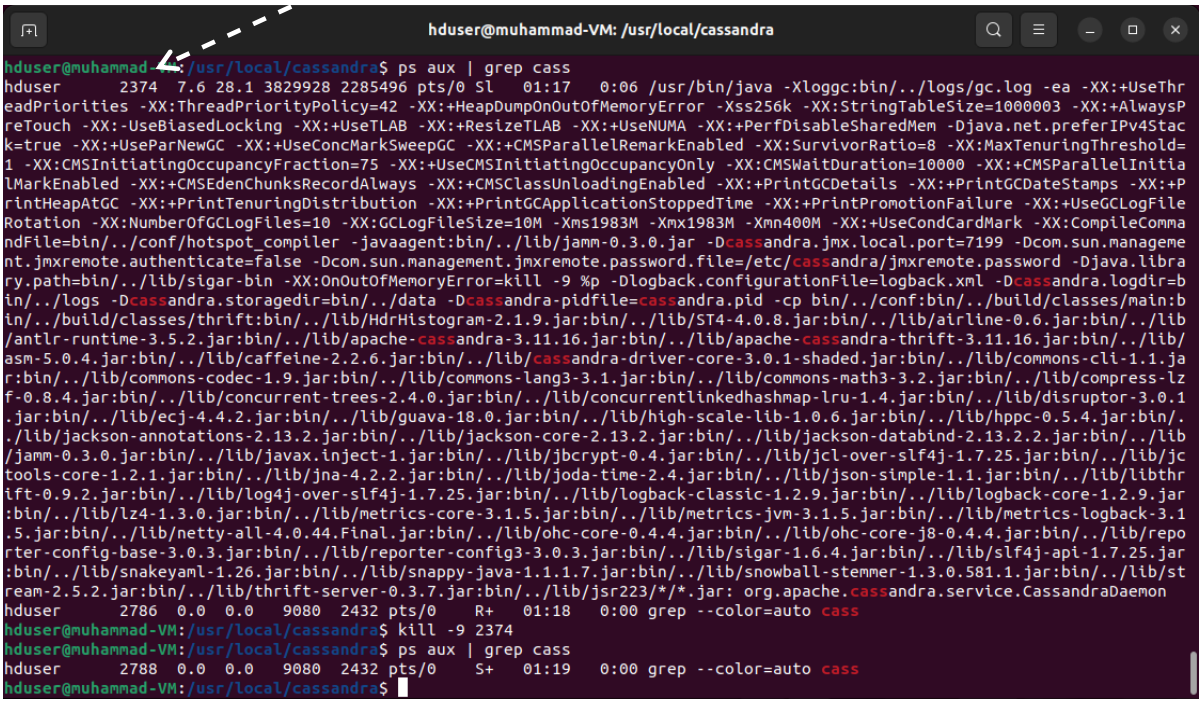

This store the PID of the Cassandra process in a file named `cassandra.pid` in the `local/cassandra` directory. Several messages will be dumped to the screen.

Run the following command as mentioned below

$bin/cassandra -f

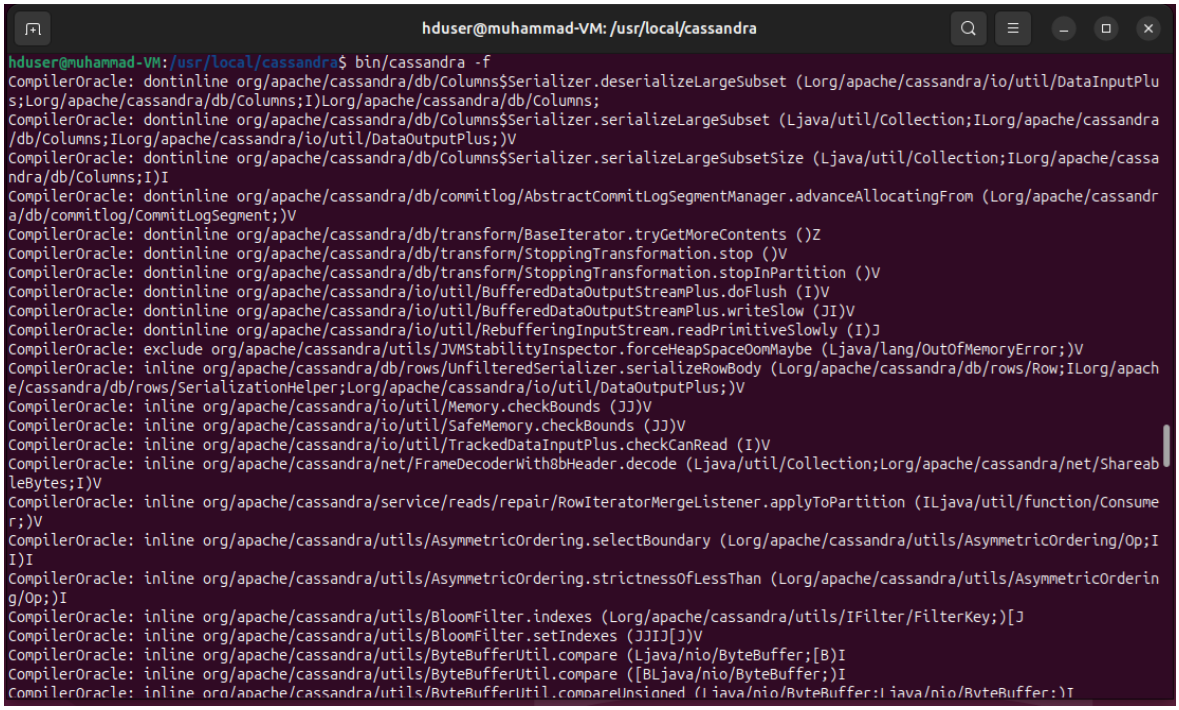

Cassandra is running smoothly. Please leave this terminal open. To check the status of Cassandra, open a new terminal by pressing Ctrl + Alt + T and run the commands shown in the screenshot below.

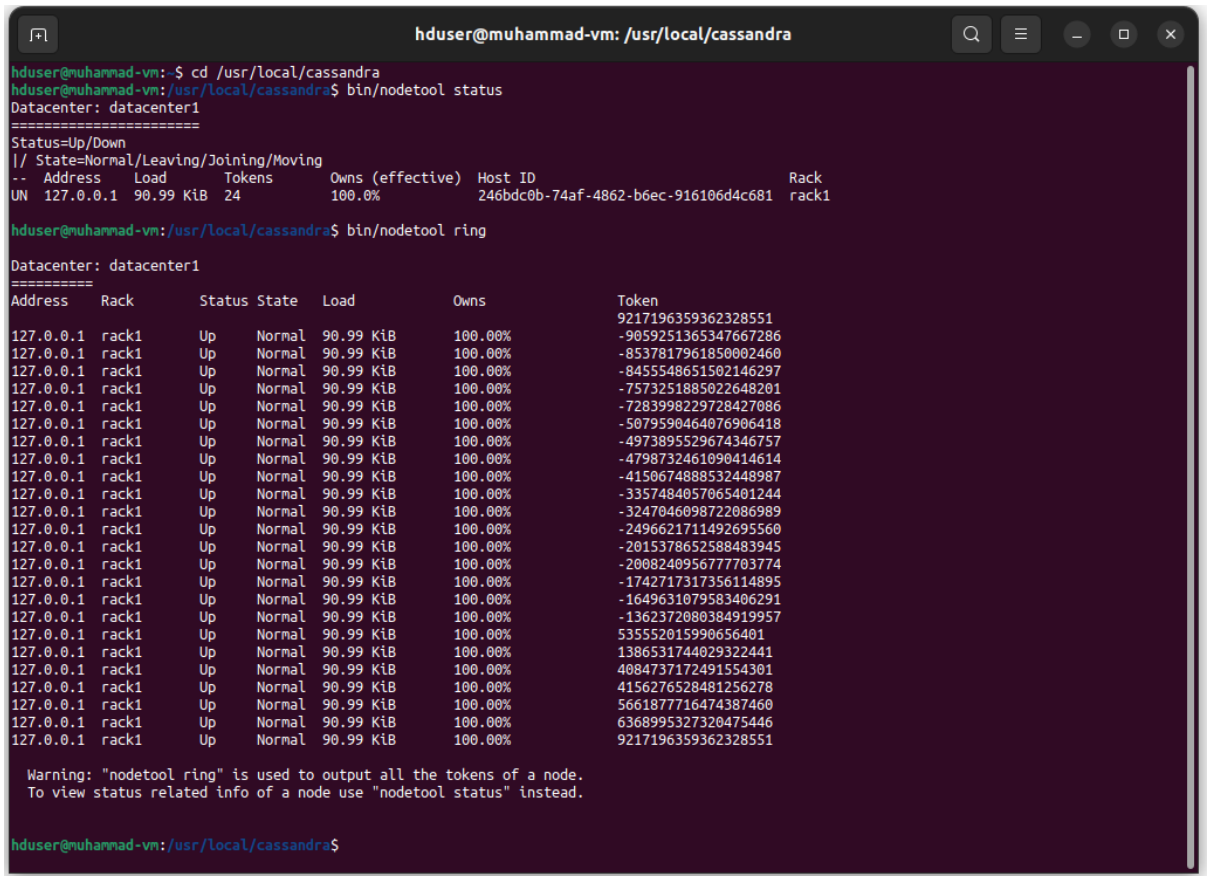

6) Install Python 2.7 on Ubuntu 24.04 LTS Linux

```
$cd /home/hduser
```

```
$sudo apt update
```

Install dependencies for Python 2.7

```
$sudo apt install -y build-essential checkinstall libncursesw5-dev libssl-dev libsqlite3-dev tk-dev libgdbm-dev libc6-dev libbz2-dev libffi-dev
```

Download the Python 2.7 Source code

```
$wget https://www.python.org/ftp/python/2.7.18/Python-2.7.18.tgz
$tar -xvf Python-2.7.18.tgz
```

Compile and Install Python 2.7 on Ubuntu 24.04

```
$cd Python-2.7.18
$./configure --enable-optimizations
$sudo make install
$python2 -V
```

A quick introduction to the data model

Now that we have a Cassandra cluster running on our local machine, we will demonstrate its use with some quick examples. We will start with **`cqlsh`**, and use that as our primary means of working with the Cassandra data model.

7) Using Cassandra with **cqlsh**: To start working with Cassandra, let's start the **Cassandra Query Language** (**CQL**) shell. The shell interface will allow us to execute CQL commands to define, query, and modify our data. As this is a new cluster and we have turned on authentication and authorization, we will use the default cassandra and cassandra username and password, as follows:

`$cd /usr/local/cassandra/`

`$bin/cqlsh`

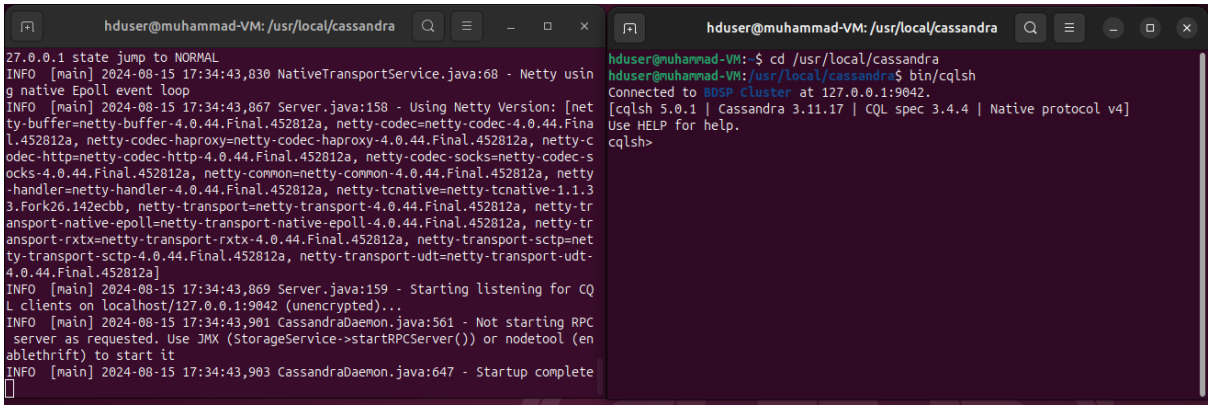

One terminal showed that the Cassandra is running, and you can execute the Cassandra database commands on the other terminal.

`cassandra@cqlsh> describe cluster;`

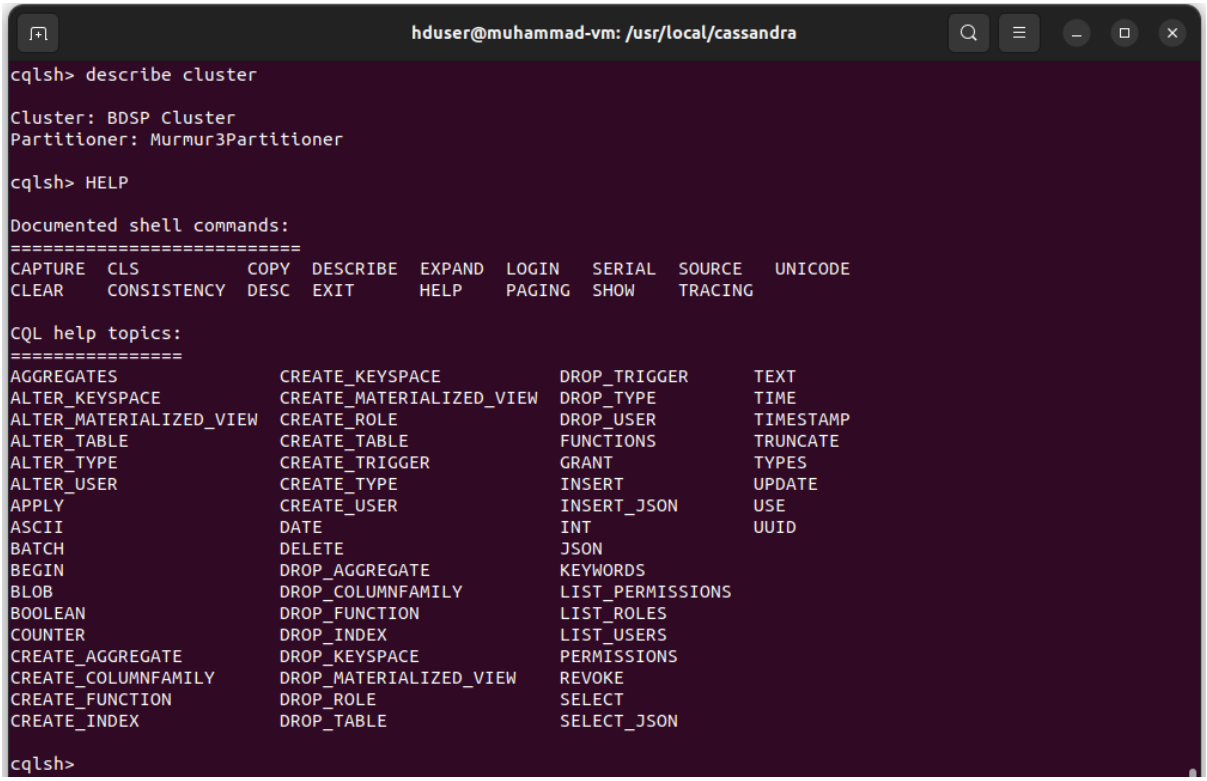

cqlsh>DESCRIBE KEYSPACES;

Check all the tables that are defined in the keyspace.

cqlsh>DESCRIBE KEYSPACE system;

cqlsh>CREATE KEYSPACE vehicle_tracker WITH REPLICATION = { 'class' : 'SimpleStrategy', 'replication_factor' : 1 };

cqlsh>DESCRIBE KEYSPACES;

Check the screenshot of this command on the next page of the tutorial. If you like to drop the keyspace

cqlsh>DROP KEYSPACE vehicle_tracker;

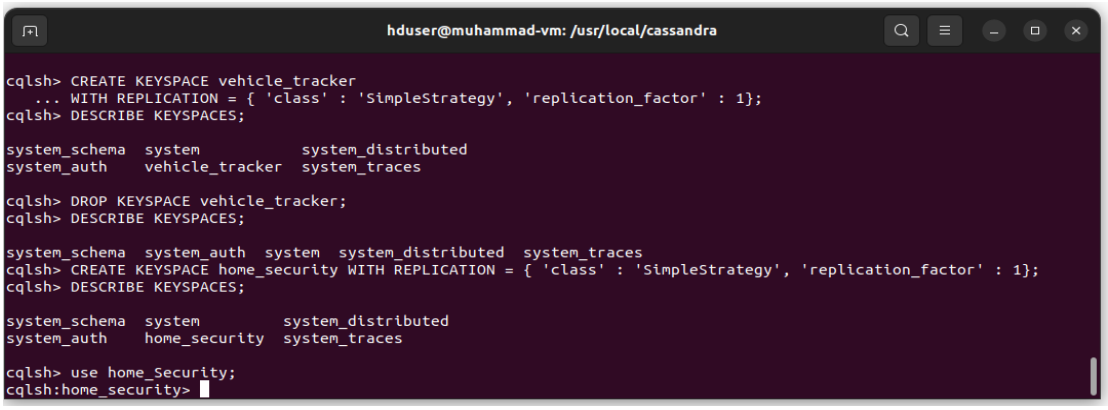

If you would like to know the details of the commands, please check the website: https://cassandra.apache.org/doc/latest/cassandra/developing/cql/ddl.html

cqlsh>USE home_Security;

Follow the screenshot to create the Table in the collection **'home_security'**

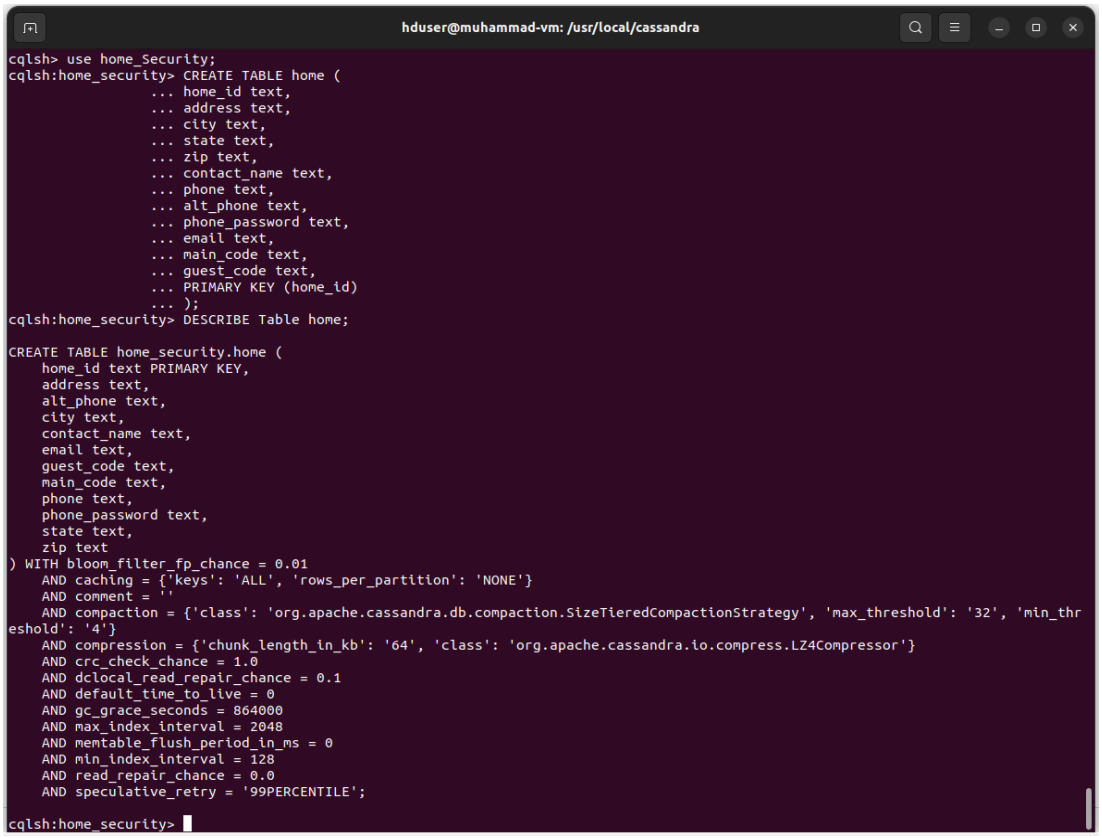

Create another table named as ‘activity’ inside the collection ‘home_security’ and the screenshots are mentioned below

8) Copy the data from **csv** file. Download the file **‘events.csv’** and **‘homes.csv’** from Moodle in the **‘Downloads’** folder on VM (This is not Hadoop and it is your local Ubuntu machine) and write the command as mentioned in the screenshot.

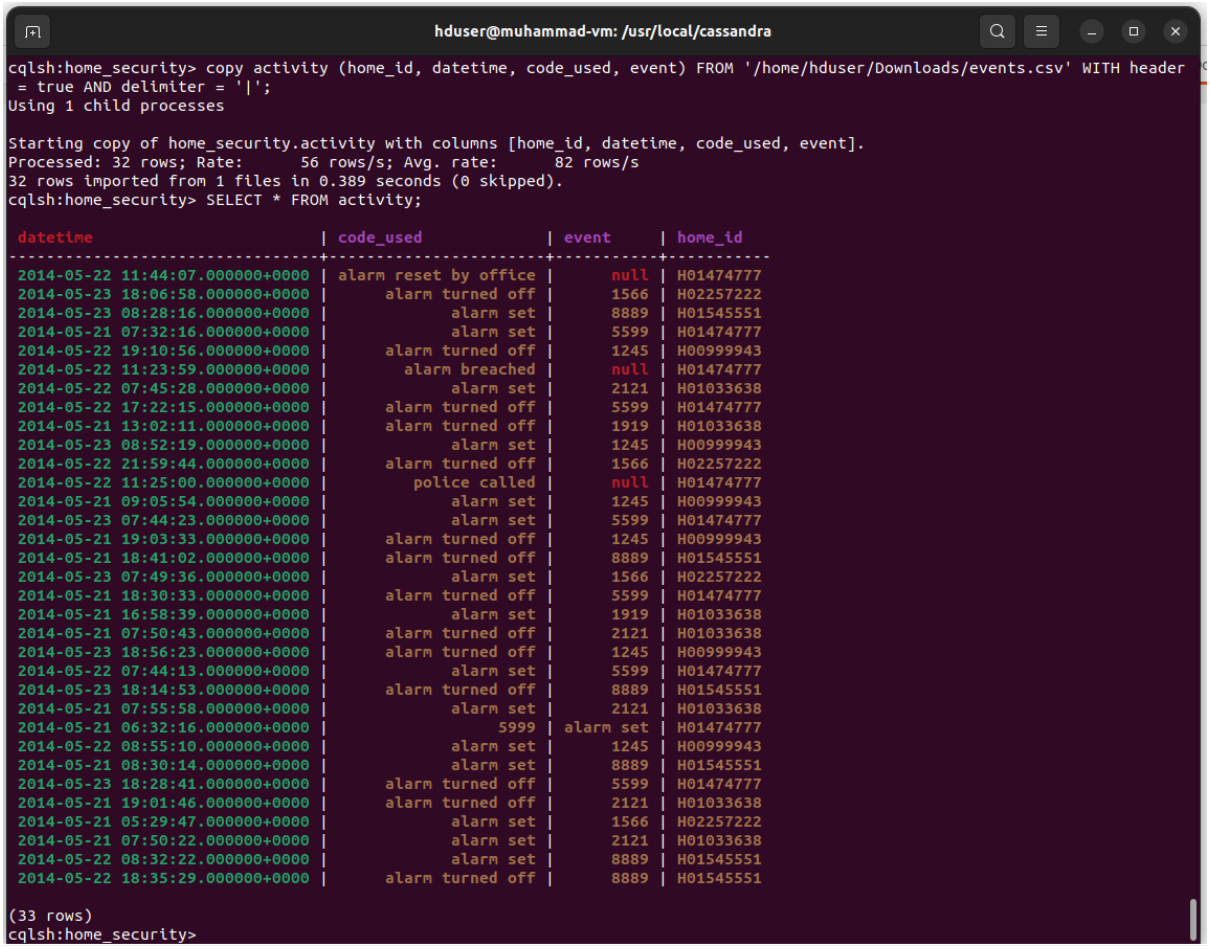

```
hduser@muhammad-vm: /usr/local/cassandra

cqlsh:home_security> copy activity (home_id, datetime, code_used, event) FROM '/home/hduser/Downloads/events.csv' WITH header
 = true AND delimiter = '|';
Using 1 child processes

Starting copy of home_security.activity with columns [home_id, datetime, code_used, event].
Processed: 32 rows; Rate:      56 rows/s; Avg. rate:      82 rows/s
32 rows imported from 1 files in 0.389 seconds (0 skipped).
cqlsh:home_security> SELECT * FROM activity;

 datetime                     | code_used            | event     | home_id
------------------------------+----------------------+-----------+-----------
 2014-05-22 11:44:07.000000+0000 | alarm reset by office |      null | H01474777
 2014-05-23 18:06:58.000000+0000 |      alarm turned off |      1566 | H02257222
 2014-05-23 08:28:16.000000+0000 |             alarm set |      8889 | H01545551
 2014-05-21 07:32:16.000000+0000 |             alarm set |      5599 | H01474777
 2014-05-22 19:10:56.000000+0000 |      alarm turned off |      1245 | H00999943
 2014-05-22 11:23:59.000000+0000 |        alarm breached |      null | H01474777
 2014-05-22 07:45:28.000000+0000 |             alarm set |      2121 | H01033638
 2014-05-22 17:22:15.000000+0000 |      alarm turned off |      5599 | H01474777
 2014-05-21 13:02:11.000000+0000 |      alarm turned off |      1919 | H01033638
 2014-05-23 08:52:19.000000+0000 |             alarm set |      1245 | H00999943
 2014-05-22 21:59:44.000000+0000 |      alarm turned off |      1566 | H02257222
 2014-05-22 11:25:00.000000+0000 |         police called |      null | H01474777
 2014-05-21 09:05:54.000000+0000 |             alarm set |      1245 | H00999943
 2014-05-23 07:44:23.000000+0000 |             alarm set |      5599 | H01474777
 2014-05-21 19:03:33.000000+0000 |      alarm turned off |      1245 | H00999943
 2014-05-21 18:41:02.000000+0000 |      alarm turned off |      8889 | H01545551
 2014-05-23 07:49:36.000000+0000 |             alarm set |      1566 | H02257222
 2014-05-21 18:30:33.000000+0000 |      alarm turned off |      5599 | H01474777
 2014-05-21 16:58:39.000000+0000 |             alarm set |      1919 | H01033638
 2014-05-21 07:50:43.000000+0000 |      alarm turned off |      2121 | H01033638
 2014-05-23 18:56:23.000000+0000 |      alarm turned off |      1245 | H00999943
 2014-05-22 07:44:13.000000+0000 |             alarm set |      5599 | H01474777
 2014-05-23 18:14:53.000000+0000 |      alarm turned off |      8889 | H01545551
 2014-05-21 07:55:58.000000+0000 |             alarm set |      2121 | H01033638
 2014-05-21 06:32:16.000000+0000 |                  5999 | alarm set | H01474777
 2014-05-22 08:55:10.000000+0000 |             alarm set |      1245 | H00999943
 2014-05-21 08:30:14.000000+0000 |             alarm set |      8889 | H01545551
 2014-05-23 18:28:41.000000+0000 |      alarm turned off |      5599 | H01474777
 2014-05-21 19:01:46.000000+0000 |      alarm turned off |      2121 | H01033638
 2014-05-21 05:29:47.000000+0000 |             alarm set |      1566 | H02257222
 2014-05-21 07:50:22.000000+0000 |             alarm set |      2121 | H01033638
 2014-05-22 08:32:22.000000+0000 |             alarm set |      8889 | H01545551
 2014-05-22 18:35:29.000000+0000 |      alarm turned off |      8889 | H01545551

(33 rows)
cqlsh:home_security>
```

```
hduser@muhammad-vm: /usr/local/cassandra

 2014-05-22 08:32:22.000000+0000 |             alarm set |      8889 | H01545551
 2014-05-22 18:35:29.000000+0000 |      alarm turned off |      8889 | H01545551

(33 rows)
cqlsh:home_security> copy home (home_id, address, city, state, zip, contact_name, phone, alt_phone, phone_password, email,
main_code, guest_code) FROM '/home/hduser/Downloads/homes.csv' WITH header = true AND delimiter = '|';
Using 1 child processes

Starting copy of home_security.home with columns [home_id, address, city, state, zip, contact_name, phone, alt_phone, phone
_password, email, main_code, guest_code].
Processed: 5 rows; Rate:      10 rows/s; Avg. rate:      14 rows/s
5 rows imported from 1 files in 0.355 seconds (0 skipped).
cqlsh:home_security> SELECT * FROM home;

 home_id   | address                  | alt_phone    | city          | contact_name      | email                | guest_co
de | main_code | phone        | phone_password | state | zip
-----------+--------------------------+--------------+---------------+-------------------+----------------------+---------
---+-----------+--------------+----------------+-------+-------
 H01474777 |       518 Crestview Drive |         null | Beverly Hills |      Jed Clampett |    jclampett@bhb.com |       77
78 |      5599 | 310-775-4011 |            oil |    CA | 90046
 H01033638 |      129 West 81st Street | 212-483-1072 |      New York |    Jerry Seinfeld |   jerry@seinfeld.com |       19
19 |      2121 | 212-440-5657 |         elaine |    NY | 10024
 H02257222 | 1164 Morning Glory Circle |         null |      Westport |   Darrin Stephens | darrin@bewitched.com |       25
66 |      1566 | 860-777-7564 |            sam |    CT | 06880
 H01545551 |   565 North Clinton Drive |         null |     Milwaukee | Arthur Fonzarelli | fonzie@happydays.com |       44
45 |      8889 | 414-858-1033 |         harley |    WI | 53525
 H00999943 |      245 East 73rd Street | 212-495-5755 |      New York |   Carrie Bradshaw |      carrie@sitc.com |       12
56 |      1245 | 212-520-0250 |          prada |    NY | 10021

(5 rows)
cqlsh:home_security>
```

9) Export the data from the Cassandra table to ‘csv’ file on your local Ubuntu machine.

```
cqlsh:home_security> COPY home_security.activity(home_id, datetime, code_used, event) TO '/home/hduser/Downloads/activity_
from_cassandra.csv' WITH header = true AND delimiter = '|';
Using 1 child processes

Starting copy of home_security.activity with columns [home_id, datetime, code_used, event].
Processed: 33 rows; Rate:     385 rows/s; Avg. rate:     385 rows/s
33 rows exported to 1 files in 0.101 seconds.
cqlsh:home_security> exit
hduser@muhammad-vm:/usr/local/cassandra$
```

The output file will be stored in 'Downloads' folder as shown below on Ubuntu VM. You might see some other files in the Download folder than this screenshot. Make sure that activity_from_cassandra.csv must be present.

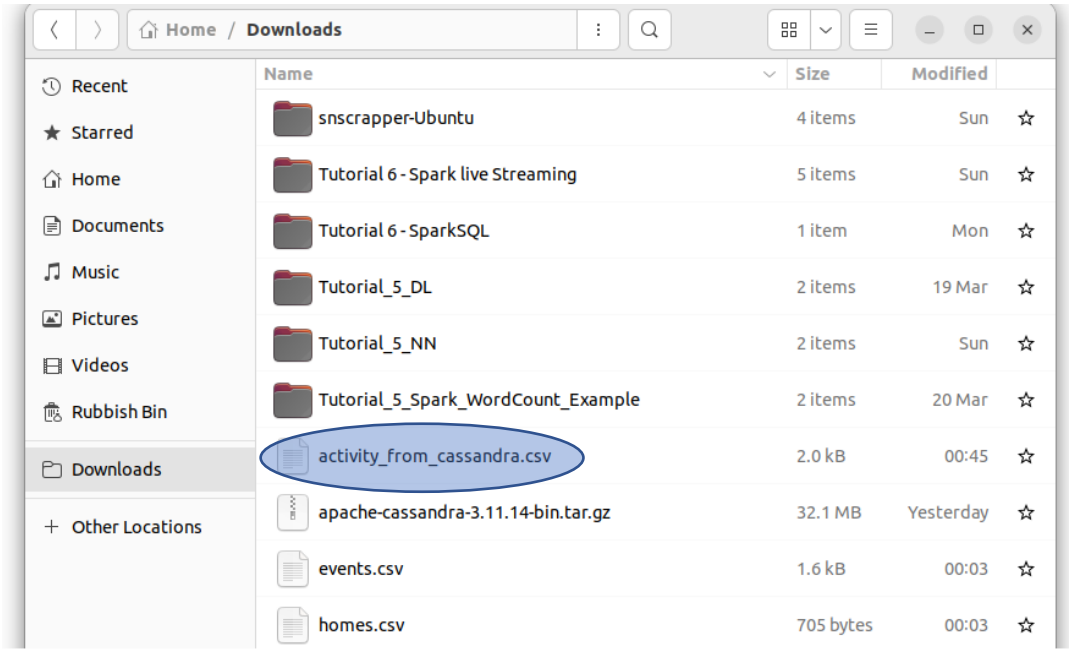

Steps to load data file from local ubuntu system to Cassandra Table

1) Create a file named as "employees_data.csv" and insert the records as mentioned below

$nano employees_data.csv

employee_id,firstname,lastname,department,city
1,Peter,Mark,Engineering,Dublin
2,Sean,Kelly,Physics,Dublin
3,Derek,Monahan,IT,Galway
4,Miles,Turner,Medical,Cork
5,Sarah,Hayes,Nursing,Cork

Or download the file "employees_data.csv" from Moodle.

2) Follow the sequence of commands to load data into Table "employees_data" and Keyspace named as "employees".

$cd /usr/local/cassandra

hduser@muhammad-VM:/usr/local/cassandra$ bin/cqlsh

Connected to BDSP Cluster at 127.0.0.1:9042.

cqlsh> CREATE KEYSPACE employees WITH replication = {'class':'

```
SimpleStrategy', 'replication_factor' : 1};
cqlsh> CREATE TABLE employees.employees_data (employee_id int PRI-
MARY KEY, firstname text, lastname text, department text, city
text);
cqlsh> USE employees;
cqlsh:employees> COPY employees.employees_data (employee_id, first-
name, lastname, department, city) FROM '/home/hduser/Downloads/
employees_data.csv' WITH HEADER = true;
cqlsh:employees> select * from employees.employees_data;
```

3) The screenshot showed the sequence of commands as shown below.

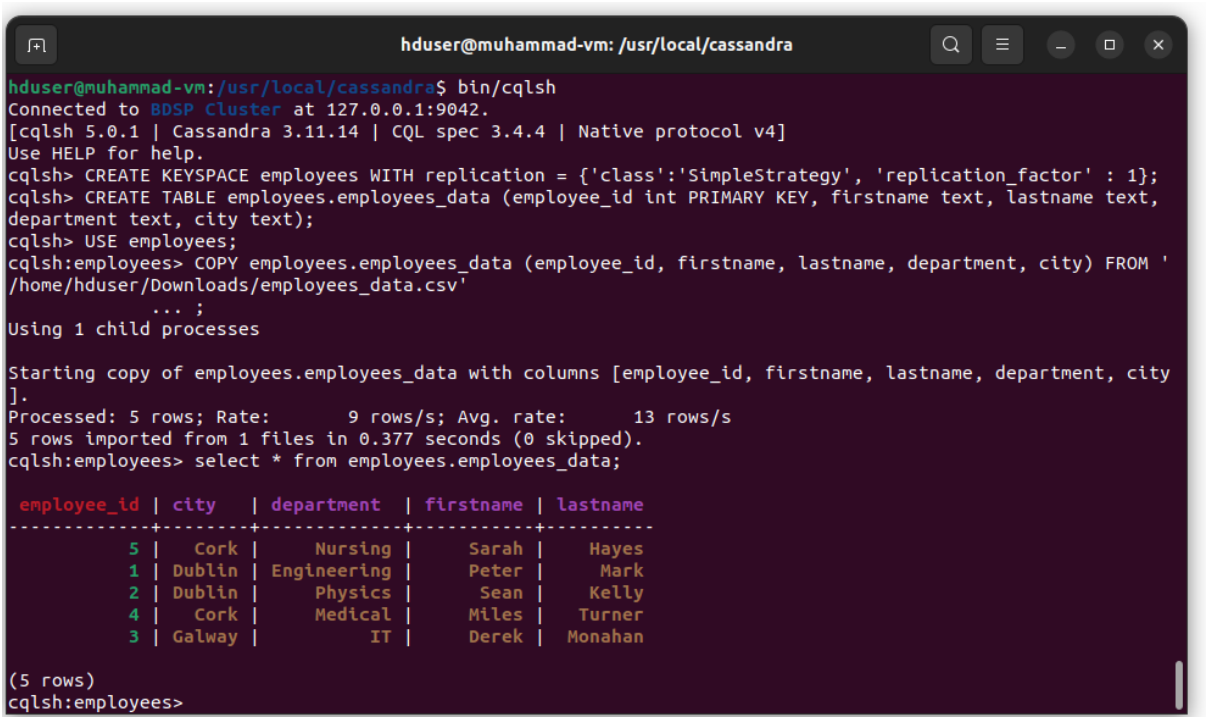

4) Learn from the book reference provided in references for further understanding of Cassandra query language and perform queries on the datasets of your choice.

References:

- https://cassandra.apache.org/doc/latest/cassandra/cql/ddl.html
- Cassandra: The Definitive Guide, (Revised) Third Edition, 3rd Edition, Jeff Carpenter, Eben Hewitt, O'Reilly Media, Inc., January 2022.
- Installation instructions: https://www.cloudduggu.com/cassandra/installation/
- https://linux.how2shout.com/how-to-install-python-2-7-on-ubuntu-24-04-noble-lts-linux/